

CAPABL  
SUMMER INTERNSHIP PROGRAM -DATA-SCIENCE

MACHINE LEARNING PROJECT REPORT  
PREDICTING TAXI TRIP DURATION



Submitted by -  
Tejas Prasad

# STUDYING THE DATA SET

For our project we are provided with 2 datasets, 1 for training the model and another one for testing it. Hence we do not need to divide our datasets any further. Let us briefly analyse our datasets.

## Training data

The training dataset contains 1,458,644 rows and 11 columns. Each row represents a taxi trip with information such as a unique identifier for the trip, the vendor or taxi company associated with the trip, the pickup and dropoff date and time, the number of passengers, longitude and latitude coordinates for the pickup and dropoff locations, a flag indicating if the trip data was stored before forwarding, and the duration of the trip in seconds.

## Testing data

The testing dataset consists of 625,134 rows and 9 columns. Each row represents a taxi trip with information such as a unique identifier for the trip, the vendor or taxi company associated with the trip, the pickup date and time, the number of passengers, longitude and latitude coordinates for the pickup and dropoff locations, and a flag indicating if the trip data was stored before forwarding. The dataset includes numerical and categorical columns. The pickup\_datetime column requires datetime processing for further analysis.

Both the datasets include both numerical and categorical columns, and some columns may require preprocessing, such as converting datetime objects, handling missing values, and potentially correcting data types. Before performing any analysis or building predictive models, it is essential to ensure data quality and conduct necessary data preprocessing steps.

# IMPORTING THE NECESSARY LIBRARIES

Before I begin the project, i need to import the necessary libraries in python -

```
[59] import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
from sklearn.linear_model import LinearRegression
from geopy.distance import geodesic
```

The libraries and modules ive imported here are -

1. `pandas` (imported as `pd`): A powerful data manipulation library that provides data structures and functions for working with structured data, like DataFrames.
2. `numpy` (imported as `np`): A fundamental library for numerical computations in Python. It provides support for large, multi-dimensional arrays and matrices, along with a collection of mathematical functions.
3. `matplotlib.pyplot` (imported as `plt`): A popular plotting library used to create various visualizations, such as line plots, scatter plots, bar plots, and more. It works seamlessly with NumPy arrays and Pandas DataFrames.
4. `sklearn.metrics.mean\_squared\_error`: A function from scikit-learn (imported as `mean\_squared\_error`) used to calculate the mean squared error, a common metric for regression tasks, to evaluate the performance of a regression model by comparing the true target values with the predicted values.
5. `sklearn.metrics.mean\_absolute\_error`: A function from scikit-learn (imported as `mean\_absolute\_error`) used to calculate the mean absolute error, another regression metric, which measures the average absolute difference between the true and predicted values.
6. `sklearn.metrics.r2\_score`: A function from scikit-learn (imported as `r2\_score`) used to calculate the R-squared score (coefficient of determination), which indicates how well the regression model fits the observed data.
7. `sklearn.linear\_model.LinearRegression`: A class from scikit-learn (imported as `LinearRegression`) used to implement linear regression models, one of the simplest and most widely used regression techniques for predicting continuous target variables based on the input features.
8. `geopy.distance.geodesic`: A function from the geopy library used for calculating geodesic distances between two points given their latitude and longitude coordinates. This can be helpful for calculating distances between pickup and dropoff locations in geographical applications.

# INITIAL DATASET DESCRIPTION

Here's what the dataset initially looks like. I have used the pandas csv reader function to read the training dataset into 'data' and testing dataset into 'data1'. I have also used the head function to print the first few rows of each feature -

```
[4] data=pd.read_csv("train.csv")
print(data.head())
```

	id	vendor_id	pickup_datetime	dropoff_datetime	\
0	id2875421	2	2016-03-14 17:24:55	2016-03-14 17:32:30	
1	id2377394	1	2016-06-12 00:43:35	2016-06-12 00:54:38	
2	id3858529	2	2016-01-19 11:35:24	2016-01-19 12:10:48	
3	id3504673	2	2016-04-06 19:32:31	2016-04-06 19:39:40	
4	id2181028	2	2016-03-26 13:30:55	2016-03-26 13:38:10	

	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	\
0	1.0	-73.982155	40.767937	-73.964630	
1	1.0	-73.980415	40.738564	-73.999481	
2	1.0	-73.979027	40.763939	-74.005333	
3	1.0	-74.010040	40.719971	-74.012268	
4	1.0	-73.973053	40.793209	-73.972923	

	dropoff_latitude	store_and_fwd_flag	trip_duration
0	40.765602	N	455.0
1	40.731152	N	663.0
2	40.710087	N	2124.0
3	40.706718	N	429.0
4	40.782520	N	435.0

```
[24] data1=pd.read_csv("test.csv")
print(data1.head())
```

	id	vendor_id	pickup_datetime	dropoff_datetime	\
0	id2875421	2	2016-03-14 17:24:55	2016-03-14 17:32:30	
1	id2377394	1	2016-06-12 00:43:35	2016-06-12 00:54:38	
2	id3858529	2	2016-01-19 11:35:24	2016-01-19 12:10:48	
3	id3504673	2	2016-04-06 19:32:31	2016-04-06 19:39:40	
4	id2181028	2	2016-03-26 13:30:55	2016-03-26 13:38:10	

	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	\
0	1.0	-73.982155	40.767937	-73.964630	
1	1.0	-73.980415	40.738564	-73.999481	
2	1.0	-73.979027	40.763939	-74.005333	
3	1.0	-74.010040	40.719971	-74.012268	
4	1.0	-73.973053	40.793209	-73.972923	

	dropoff_latitude	trip_duration	distance_km
0	40.765602	455.0	1.502172
1	40.731152	663.0	1.808660
2	40.710087	2124.0	6.379687
3	40.706718	429.0	1.483632
4	40.782520	435.0	1.187038

I have further used the info() function to get a more technical description of the datasets with information about each feature -

```
[9] data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1311508 entries, 0 to 1311507
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    1311508 non-null  object
1   vendor_id            1311508 non-null  int64
2   pickup_datetime      1311508 non-null  object
3   dropoff_datetime     1311508 non-null  object
4   passenger_count      1311508 non-null  float64
5   pickup_longitude     1311508 non-null  float64
6   pickup_latitude     1311508 non-null  float64
7   dropoff_longitude    1311508 non-null  float64
8   dropoff_latitude     1311508 non-null  float64
9   store_and_fwd_flag   1311508 non-null  object
10  trip_duration        1311508 non-null  float64
dtypes: float64(6), int64(1), object(4)
memory usage: 120.1+ MB
```

```
[31] data1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 625134 entries, 0 to 625133
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    625134 non-null  object
1   vendor_id            625134 non-null  int64
2   pickup_datetime      625134 non-null  object
3   passenger_count      625134 non-null  int64
4   pickup_longitude     625134 non-null  float64
5   pickup_latitude     625134 non-null  float64
6   dropoff_longitude    625134 non-null  float64
7   dropoff_latitude     625134 non-null  float64
8   store_and_fwd_flag   625134 non-null  object
dtypes: float64(4), int64(2), object(3)
memory usage: 42.9+ MB
```

# DATA PRE-PROCESSING

In order to process the datasets given and make it more usable, i have taken multiple measures like checking for missing and NULL values, extracting new features from existing ones and so on. Let me walk you through the steps ive taken.

## Checking for NULL and missing values

By using the isnull() function on both the testing and training datasets we can see if there are and null or missing values

✓ 1s	[6] data.isnull().sum()	[9] data1.isnull().sum()
	id 0	id 0
	vendor_id 0	vendor_id 0
	pickup_datetime 0	pickup_datetime 0
	dropoff_datetime 0	dropoff_datetime 0
	passenger_count 1	passenger_count 0
	pickup_longitude 1	pickup_longitude 0
	pickup_latitude 1	pickup_latitude 0
	dropoff_longitude 1	dropoff_longitude 0
	dropoff_latitude 1	dropoff_latitude 0
	store_and_fwd_flag 1	store_and_fwd_flag 0
	trip_duration 1	dtype: int64
	dtype: int64	

Here we can see that our testing data has no null values but our training data does, so ill go ahead and remove them using the drop function

[10] data=data.dropna() data.isnull().sum()
id 0
vendor_id 0
pickup_datetime 0
dropoff_datetime 0
passenger_count 0
pickup_longitude 0
pickup_latitude 0
dropoff_longitude 0
dropoff_latitude 0
store_and_fwd_flag 0
trip_duration 0
dtype: int64

Now we can see that all the null and missing values have been removed.

## Feature extraction

In our data set we have been given 4 features, The pickup and drop off latitudes and longitudes. Using these co-ordinates we can extract a more meaningful feature which will give is the distance of each of the rides. We can do this for both the datasets

```
[15] pickup_coords = list(zip(data['pickup_latitude'], data['pickup_longitude']))
      dropoff_coords = list(zip(data['dropoff_latitude'], data['dropoff_longitude']))

      data['distance_km'] = [geodesic(pickup, dropoff).kilometers for pickup, dropoff in zip(pickup_coords, dropoff_coords)]

[65] pickup_coords = list(zip(data1['pickup_latitude'], data1['pickup_longitude']))
      dropoff_coords = list(zip(data1['dropoff_latitude'], data1['dropoff_longitude']))

      data1['distance_km'] = [geodesic(pickup, dropoff).kilometers for pickup, dropoff in zip(pickup_coords, dropoff_coords)]
```

## Dropping Obsolete Features

In our datasets we have a feature called “store\_and\_fwd\_flag” which basically indicates if the details of a ride have been forwarded to the operator. For this feature, if most of the rides have a flag ‘Yes’ , then this could potentially have an impact on our model. However if most of the flags indicate ‘no’ then this feature could be useless to our model. I will go ahead and check the ratio of yes and no flags

```
[11] print(data['store_and_fwd_flag'].value_counts())

N    1304277
Y      7231
Name: store_and_fwd_flag, dtype: int64
```

Here I can see that there are insignificantly less number of yes flags. Hence i will go ahead and drop this feature.

```
data.drop('store_and_fwd_flag', axis=1, inplace=True)
data1.drop('store_and_fwd_flag', axis=1, inplace=True)
```

## Checking for outliers

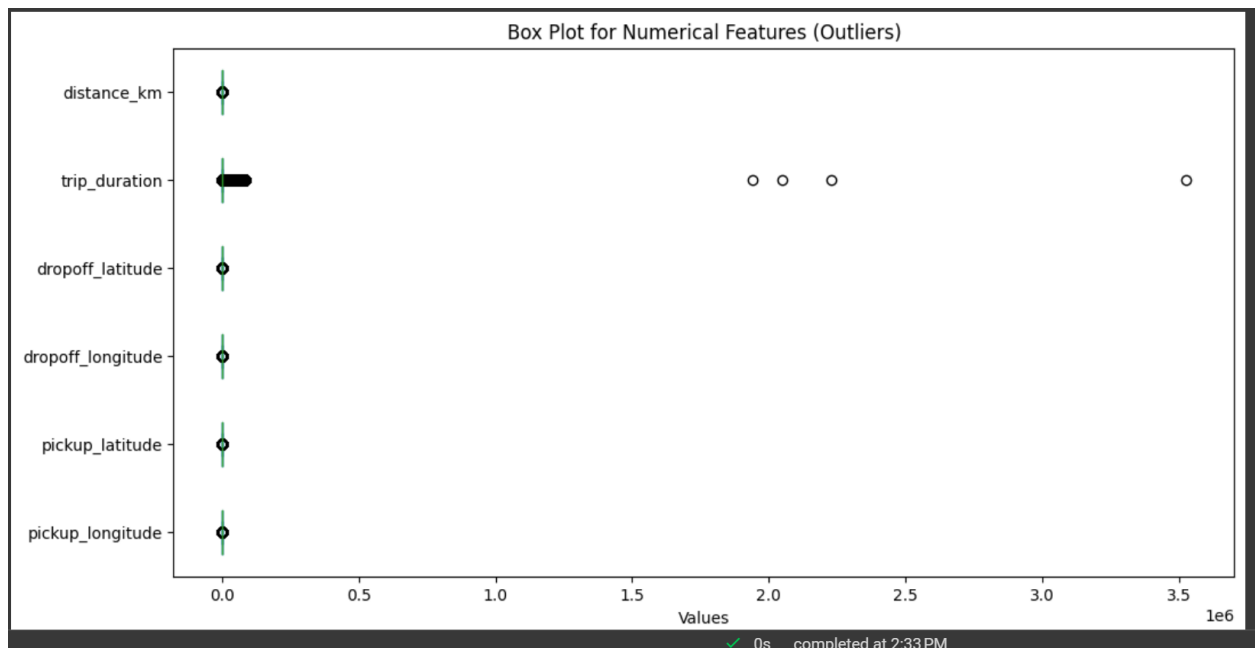
I will now go ahead and check for any outliers in our dataset. For this ill use the matplotlib library to plot a boxplot for the given dataset

```

+ Code + Text
numerical_columns = ['pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude', 'trip_duration', 'distance_km']
data[numerical_columns].plot(kind='box', vert=False, figsize=(12, 6))
plt.title("Box Plot for Numerical Features (Outliers)")
plt.xlabel("Values")
plt.show()

```

I will get a graph that looks something like this



From this i can see that there arent any significant outliers so i can move on from here

## Description of Pre-Processed data

After i have pre-processed my datasets it looks like this

```

data.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1311508 entries, 0 to 1311507
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   id                     1311508 non-null object
1   vendor_id              1311508 non-null int64
2   pickup_datetime        1311508 non-null object
3   dropoff_datetime       1311508 non-null object
4   passenger_count        1311508 non-null float64
5   pickup_longitude       1311508 non-null float64
6   pickup_latitude        1311508 non-null float64
7   dropoff_longitude      1311508 non-null float64
8   dropoff_latitude       1311508 non-null float64
9   distance_km            1311508 non-null float64
dtypes: float64(6), int64(1), object(3)
memory usage: 142.3+ MB

[67] data1.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 625134 entries, 0 to 625133
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   id                     625134 non-null object
1   vendor_id              625134 non-null int64
2   pickup_datetime        625134 non-null object
3   passenger_count        625134 non-null int64
4   pickup_longitude       625134 non-null float64
5   pickup_latitude        625134 non-null float64
6   dropoff_longitude      625134 non-null float64
7   dropoff_latitude       625134 non-null float64
8   distance_km            625134 non-null float64
dtypes: float64(5), int64(2), object(2)
memory usage: 42.9+ MB

```

Here you can see that the data now has a feature 'distance\_km' but does not have the 'store\_and\_fwd\_flag'. I would call this a success but I have made a mistake while dropping the features and accidentally removed the time\_duration feature from the training dataset which is very crucial to us. Hence i will create another variable data3 which has concatenation of data and the time\_duration feature from the original dataset.

```
data3 = pd.concat([data3, data['time_duration']], axis=1)
```

The data3 looks something like this

```
[64]
data3.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1220005 entries, 0 to 1220004
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    1220005 non-null object
1   vendor_id             1220005 non-null int64
2   pickup_datetime       1220005 non-null object
3   dropoff_datetime      1220005 non-null object
4   passenger_count       1220005 non-null float64
5   pickup_longitude      1220005 non-null float64
6   pickup_latitude       1220005 non-null float64
7   dropoff_longitude     1220005 non-null float64
8   dropoff_latitude      1220005 non-null float64
9   distance_km          1220005 non-null float64
10  trip_duration         1220005 non-null float64
dtypes: float64(7), int64(1), object(3)
memory usage: 111.7+ MB
```

I will be using data3 as the training set for any further applications in my model



# TRAINING AND TESTING THE MODEL

To implement a linear regression model i have used the LinearRegression module from sklearn.linear\_regression. I have trained this model on my processed training data ('data3') and tested it on my processed testing data ('data1'). The implementation looks like this

```
feature_cols = ['vendor_id', 'passenger_count', 'pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude', 'distance_km',
target_col = 'trip_duration'

# Drop rows with missing 'trip_duration' values in the training data
data3.dropna(subset=[target_col], inplace=True)

# Separate the features and target variable in training sets
X_train = data3[feature_cols]
y_train = data3[target_col]

# Create and train the linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions on the testing data
X_test = data1[feature_cols]
y_pred = model.predict(X_test)

print(y_pred)
```

```
[ 795.95450985  934.01110493  643.95412827 ... 1393.78337065 2701.77169967
1372.33865764]
```

Here, i have assigned all the feature columns to a variable featrue\_cols and the target column which is trip\_duration to the variable target\_col

I have then trained the model by assigning the target column to the y axis and the feature column to the x axis. The model will then try to fit the best line for the data points by using multiple regression

After the im done training the model have tested it on the testing data and asked it to predict the trip\_duration for the trips in the testing dataset. I the stored the predicted values in an array called y\_pred.

## RESULT AND DISCUSSION

After training my model on the training dataset, I am now able to get it to predict the trip durations for the trips in the testing dataset. I have stored it in a array `y_pred` and when i print it i get something like this

```
[69] print(y_pred)
[ 795.95450985  934.01110493  643.95412827 ... 1393.78337065 2701.77169967
 1372.33865764]
```

Since there are too many values in the testing dataset, it will not display each and every prediction.

I am not able to check the accuracy of the models predictions here since there arent any values of the trip duration given in the dataset to compare to.

However, I can still check the accuracy by extracting a feature with all the trip durations for the testing data using the pick up time and drop off time in the dataset. However the accuracy measured here will still depend on the accuracy of the time values given so i will not be doing this

In conclusion, Our model is no working and can predict the trip\_duration of any new trips given the input features